

```
from google.colab import files
uploaded = files.upload()

Choose files data2.csv
data2.csv(text/csv) - 429766 bytes, last modified: 30/10/2025 - 100% done
Saving data2.csv to data2.csv
```

```
import pandas as pd

# Load the uploaded CSV file
df = pd.read_csv("data2.csv")

# Show the first 5 rows
df.head(10)
```

	Profession	Income	Credit_card_number	Expiry	Security_code	Fraud	
0	DOCTOR	42509	3515418493460774	07/25	251	1	
1	DOCTOR	80334	213134223583196	05/32	858	1	
2	LAWYER	91552	4869615013764888	03/30	755	1	
3	LAWYER	43623	341063356109385	01/29	160	1	
4	DOCTOR	22962	4707418777543978402	11/30	102	0	
5	ENGINEER	72106	4483533221713	05/27	834	0	
6	DOCTOR	54992	348781654665997	06/30	207	0	
7	LAWYER	19996	38608157332988	12/26	433	0	
8	DOCTOR	54682	4495934016931879	01/32	872	0	
9	DOCTOR	74679	30107450695256	09/34	295	0	

Next steps: [Generate code with df](#) [New interactive sheet](#)

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10000 entries, 0 to 9999
Data columns (total 6 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   Profession      10000 non-null   object  
 1   Income          10000 non-null   int64  
 2   Credit_card_number 10000 non-null   int64  
 3   Expiry          10000 non-null   object  
 4   Security_code   10000 non-null   int64  
 5   Fraud            10000 non-null   int64  
dtypes: int64(4), object(2)
memory usage: 468.9+ KB
```

```
df.describe()
```

	Income	Credit_card_number	Security_code	Fraud	
count	10000.00000	1.000000e+04	10000.000000	10000.000000	
mean	49761.20600	3.851363e+17	863.587800	0.501600	
std	28837.72928	1.257950e+18	1484.424959	0.500022	
min	1.00000	6.040296e+10	0.000000	0.000000	
25%	24863.75000	1.800137e+14	275.000000	0.000000	
50%	49483.00000	3.512440e+15	539.500000	1.000000	
75%	74483.00000	4.594779e+15	813.250000	1.000000	
max	99986.00000	4.999697e+18	9990.000000	1.000000	

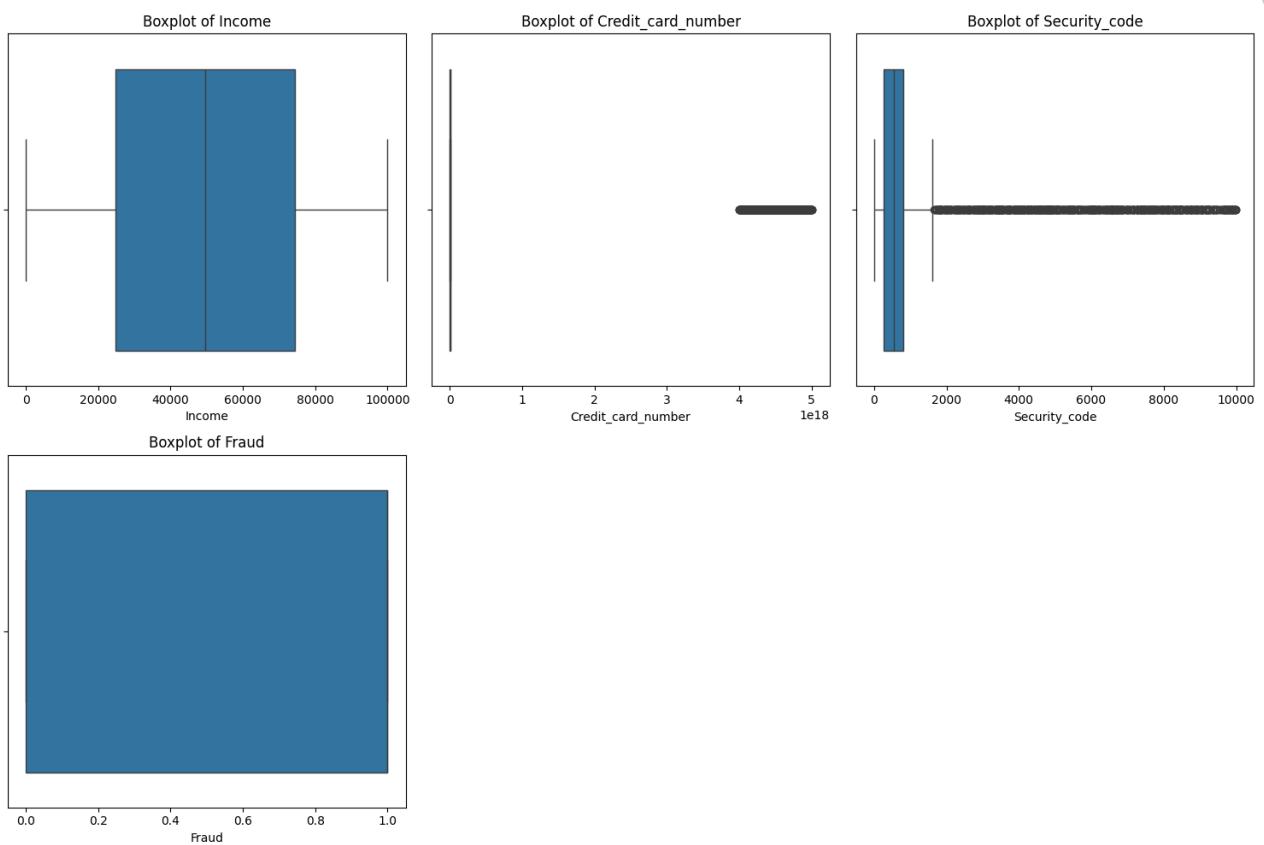
```
df.dtypes
```

```
0  
Profession    object  
Income        int64  
Credit_card_number  int64  
Expiry        object  
Security_code   int64  
Fraud          int64  
  
dtype: object
```

```
# To know is there any missing value in the dataset.  
df.isnull().sum()
```

```
0  
Profession    0  
Income        0  
Credit_card_number  0  
Expiry        0  
Security_code   0  
Fraud          0  
  
dtype: int64
```

```
# To find out is there any outliers  
import matplotlib.pyplot as plt  
import seaborn as sns  
  
# Select numeric columns only  
numeric_cols = df.select_dtypes(include=['int64', 'float64']).columns  
  
# Plot boxplots for each numeric column  
plt.figure(figsize=(15, 10))  
for i, col in enumerate(numeric_cols, 1):  
    plt.subplot(len(numeric_cols)//3 + 1, 3, i)  
    sns.boxplot(x=df[col])  
    plt.title(f'Boxplot of {col}')  
plt.tight_layout()  
plt.show()
```



```
# Normalization of numerical Columns
from sklearn.preprocessing import MinMaxScaler

scaler = MinMaxScaler()
numeric_cols = df.select_dtypes(include=['int64', 'float64']).columns

df[numeric_cols] = scaler.fit_transform(df[numeric_cols])
df.head(10)
```

	Profession	Income	Credit_card_number	Expiry	Security_code	Fraud	⋮
0	DOCTOR	0.425144	7.031142e-04	07/25	0.025125	1.0	⠄⠄
1	DOCTOR	0.803451	4.261734e-05	05/32	0.085886	1.0	
2	LAWYER	0.915647	9.739699e-04	03/30	0.075576	1.0	
3	LAWYER	0.436285	6.820472e-05	01/29	0.016016	1.0	
4	DOCTOR	0.229644	9.415407e-01	11/30	0.010210	0.0	
5	ENGINEER	0.721158	8.846796e-07	05/27	0.083483	0.0	
6	DOCTOR	0.549992	6.974847e-05	06/30	0.020721	0.0	
7	LAWYER	0.199980	7.710018e-06	12/26	0.043343	0.0	
8	DOCTOR	0.546892	8.992292e-04	01/32	0.087287	0.0	
9	DOCTOR	0.746892	6.009773e-06	09/34	0.029530	0.0	

Next steps: [Generate code with df](#) [New interactive sheet](#)

```
# Normalization of Categorical Columns
categorical_columns = df.select_dtypes(include=['object']).columns
```

```
print("Categorical Columns: ", categorical_columns)

Categorical Columns: Index(['Profession', 'Expiry'], dtype='object')

# Convert ExpiryDate to datetime first
df['Expiry'] = pd.to_datetime(df['Expiry'], errors='coerce')

# Then check again for categorical columns
categorical_columns = df.select_dtypes(include=['object']).columns
print(categorical_columns)

Index(['Profession'], dtype='object')
```

```
data_encoded = pd.get_dummies(df, columns=categorical_columns)
data_encoded.head(10)
```

	Income	Credit_card_number	Expiry	Security_code	Fraud	Profession_DOCTOR	Profession_ENGINEER	Profession_LAWYER
0	0.425144	7.031142e-04	NaT	0.025125	1.0	True	False	False
1	0.803451	4.261734e-05	2032-05-01	0.085886	1.0	True	False	False
2	0.915647	9.739699e-04	NaT	0.075576	1.0	False	False	True
3	0.436285	6.820472e-05	NaT	0.016016	1.0	False	False	True
4	0.229644	9.415407e-01	NaT	0.010210	0.0	True	False	False
5	0.721158	8.846796e-07	NaT	0.083483	0.0	False	True	False
6	0.549992	6.974847e-05	NaT	0.020721	0.0	True	False	False
7	0.199980	7.710018e-06	NaT	0.043343	0.0	False	False	True
8	0.546892	8.992292e-04	2032-01-01	0.087287	0.0	True	False	False

Next steps: [Generate code with data_encoded](#) [New interactive sheet](#)

```
# Convert all boolean columns (True/False) to integers (1/0)
bool_cols = data_encoded.select_dtypes(include=['bool']).columns
data_encoded[bool_cols] = data_encoded[bool_cols].astype(int)
data_encoded.head(10)
```

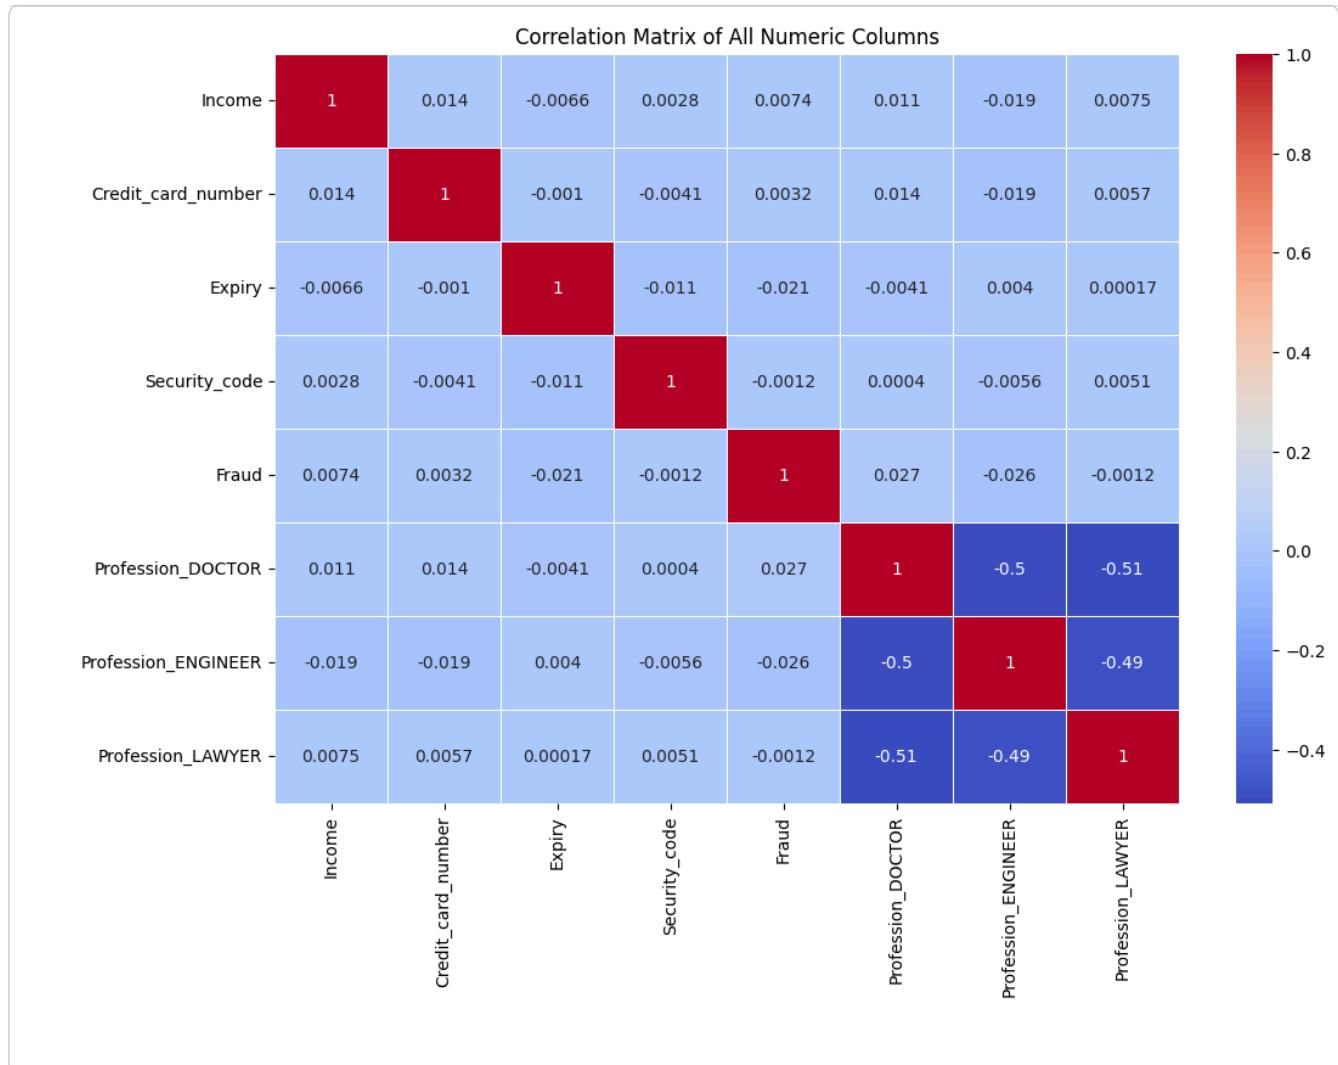
	Income	Credit_card_number	Expiry	Security_code	Fraud	Profession_DOCTOR	Profession_ENGINEER	Profession_LAWYER
0	0.425144	7.031142e-04	NaT	0.025125	1.0	1	0	0
1	0.803451	4.261734e-05	2032-05-01	0.085886	1.0	1	0	0
2	0.915647	9.739699e-04	NaT	0.075576	1.0	0	0	1
3	0.436285	6.820472e-05	NaT	0.016016	1.0	0	0	1
4	0.229644	9.415407e-01	NaT	0.010210	0.0	1	0	0
5	0.721158	8.846796e-07	NaT	0.083483	0.0	0	1	0
6	0.549992	6.974847e-05	NaT	0.020721	0.0	1	0	0
7	0.199980	7.710018e-06	NaT	0.043343	0.0	0	0	1
8	0.546892	8.992292e-04	2032-01-01	0.087287	0.0	1	0	0

Next steps: [Generate code with data_encoded](#) [New interactive sheet](#)

```
# Calculate correlation for all numeric columns
import seaborn as sns
import matplotlib.pyplot as plt

# Calculate correlation for all numeric columns
corr_matrix = data_encoded.corr()

# Plot the heatmap
plt.figure(figsize=(12, 8))
sns.heatmap(corr_matrix, annot=True, cmap='coolwarm', linewidths=0.5)
plt.title("Correlation Matrix of All Numeric Columns")
plt.show()
```



```
# Split the data into 70% training and 30% testing
from sklearn.model_selection import train_test_split

X = data_encoded.drop(columns=['Fraud', 'Expiry'])    # Feature columns
y = data_encoded['Fraud']                                # Target column

# Correct order: X_train, X_test, y_train, y_test
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)

print("Training set shape:", X_train.shape)
print("Testing set shape:", X_test.shape)

Training set shape: (7000, 6)
Testing set shape: (3000, 6)
```

```
# Decision Tree Classifier
# 1 Import the library
from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report

# 2 Create the model
dt_model = DecisionTreeClassifier(random_state=42)

# 3 Train the model on training data
dt_model.fit(X_train, y_train)

# 4 Make predictions on test data
y_pred = dt_model.predict(X_test)

# 5 Evaluate the model
print("Model Accuracy:", accuracy_score(y_test, y_pred))
print("\n Confusion Matrix:\n", confusion_matrix(y_test, y_pred))
print("\n Classification Report:\n", classification_report(y_test, y_pred))
```

Model Accuracy: 0.49333333333333335

Confusion Matrix:
[[724 752]]

[768 756]]

Classification Report:					
	precision	recall	f1-score	support	
0.0	0.49	0.49	0.49	1476	
1.0	0.50	0.50	0.50	1524	
accuracy			0.49	3000	
macro avg	0.49	0.49	0.49	3000	
weighted avg	0.49	0.49	0.49	3000	

```
#Logistic Regression
# 1 Import libraries
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report

# 2 Create the model
log_model = LogisticRegression(max_iter=1000, random_state=42)

# 3 Train the model
log_model.fit(X_train, y_train)

# 4 Predict on test data
y_pred = log_model.predict(X_test)

# 5 Evaluate performance
print(" Model Accuracy:", accuracy_score(y_test, y_pred))
print("\n Confusion Matrix:\n", confusion_matrix(y_test, y_pred))
print("\n Classification Report:\n", classification_report(y_test, y_pred))
```

✓ Model Accuracy: 0.5143333333333333

Confusion Matrix:
[[999 477]
[980 544]]

Classification Report:					
	precision	recall	f1-score	support	
0.0	0.50	0.68	0.58	1476	
1.0	0.53	0.36	0.43	1524	
accuracy			0.51	3000	
macro avg	0.52	0.52	0.50	3000	
weighted avg	0.52	0.51	0.50	3000	

```
# Random Forest Classifier
# 1) Import
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report

# 2) Create the model (solid defaults)
rf = RandomForestClassifier(
    n_estimators=300,
    max_depth=None,           # let trees grow; you can try values like 5, 10, 20
    min_samples_split=2,
    min_samples_leaf=1,
    n_jobs=-1,                # use all CPUs
    random_state=42
)

# 3) Train
rf.fit(X_train, y_train)

# 4) Predict
y_pred = rf.predict(X_test)

# 5) Evaluate
print(" Random Forest Accuracy:", accuracy_score(y_test, y_pred))
print("\n Confusion Matrix:\n", confusion_matrix(y_test, y_pred))
print("\n Classification Report:\n", classification_report(y_test, y_pred))
```

✓ Random Forest Accuracy: 0.495